

Introduction to Deep Learning

Dr. Yaohua Zang & Vincent Scholz

Professorship for Data-driven Materials Modeling
Technical University Munich

April 23, 2025

Contact Information

- **Professor:** Prof. Koutsourelakis
- **Instructor:** Dr. Yaohua Zang, Vincent Scholz
- **Email:** yaohua.zang@tum.de
- **Office Location:** ED Building, Room 0438
- **Course Website:**
<https://www.moodle.tum.de/course/view.php?id=107803>

Motivation & Applications

- **Speech Recognition**
 - Siri, Google Assistant
- **Computer Vision**
 - Image classification, object detection
- **Natural Language Processing**
 - ChatGPT, translation services
- **Scientific Applications**
 - Drug discovery, genomics, physics
- **Autonomous Vehicles**
- **Medical Diagnosis**



Figure: AI applications across industries

Foreshadowing (I)

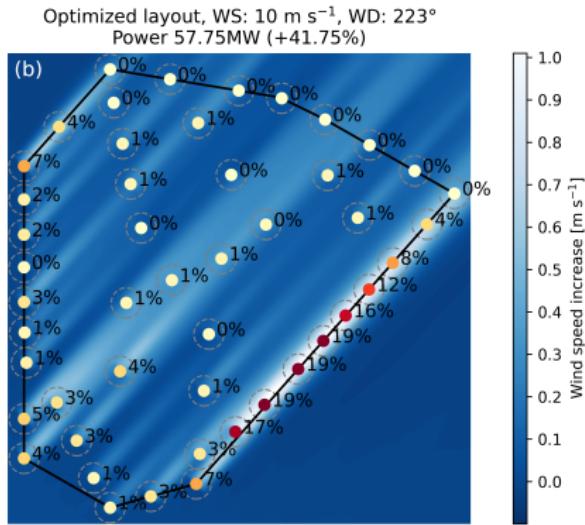
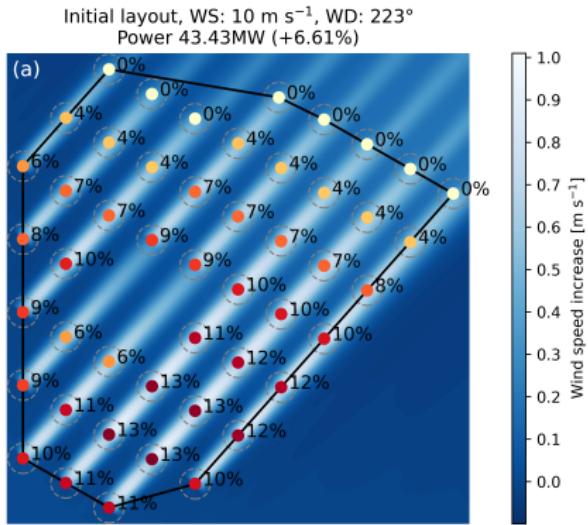


Figure: Example wind farm optimization.

Foreshadowing (II)



Figure: Example JET Fusion reactor.

Recommended Literature

- Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep Learning*. MIT Press.
- Chollet, F. (2021). *Deep Learning with Python*. Manning Publications.
- Aggarwal, C. C. (2018). *Neural Networks and Deep Learning: A Textbook*. Springer.
- Zhang, A., Lipton, Z. C., Li, M., & Smola, A. J. (2021). *Dive into Deep Learning*. Cambridge University Press.

Universal Approximation Theorem

Theorem (Universal Approximation Theorem)

A sufficiently large neural network with at least one hidden layer can approximate any continuous function to an arbitrary degree of accuracy

- **Mathematical formulation:**

- For any continuous function $u : K \rightarrow \mathbb{R}$ in a compact $K \subset \mathbb{R}^d$ and $\epsilon > 0$, there exists a MLP $NN(x; \theta)$ such that:

$$\sup_{x \in K} |u(x) - NN(x, \theta)| < \epsilon \quad (1)$$

- Explains why neural networks are universal function approximators

Table of Contents

Theorem (Universal Approximation Theorem)

A sufficiently large neural network with at least one hidden layer can approximate any continuous function to an arbitrary degree of accuracy

1 What is a neural network?

- History
- Basic Structure & Components

2 Why at least one hidden layer? (aka Deep Learning)

3 How do we make it approximate any continuous function?

- Learning tasks & Examples
- Types of Neural Networks

Table of Contents

Theorem (Universal Approximation Theorem)

A sufficiently large neural network with at least one hidden layer can approximate any continuous function to an arbitrary degree of accuracy

1 What is a neural network?

- History
- Basic Structure & Components

2 Why at least one hidden layer? (aka Deep Learning)

3 How do we make it approximate any continuous function?

- Learning tasks & Examples
- Types of Neural Networks

Artificial Neural Networks (ANNs)

- **Core of deep learning**
- Enable computers to:
 - Recognize complex patterns
 - Solve intricate problems
 - Adapt to dynamic environments
- **Revolutionary applications:**
 - Natural Language Processing
 - Self-Driving Vehicles
 - Medical Diagnosis
 - Automated Decision-Making

Evolution of Neural Networks

- **1940s-1950s:** McCulloch & Pitts - first mathematical model
- **1960s-1970s:** Rosenblatt - Perceptron (linearly separable problems)
- **1980s:** Rumelhart, Hinton, Williams - Backpropagation algorithm
- **1990s:** Applications in image recognition; "AI winter"
- **2000s:** GPU advancements, larger datasets, new architectures
- **2010s-Present:** Deep learning dominates AI
 - CNNs for image recognition
 - RNNs for sequential data
 - Transformers for NLP (ChatGPT, BERT)

Basic Structure of Neural Networks

- **Inspired by human brain**
- Each neuron performs mathematical operations
- Multiple layers of neurons
- **Components:**
 - Input Layer
 - Hidden Layers
 - Output Layer
 - Connections

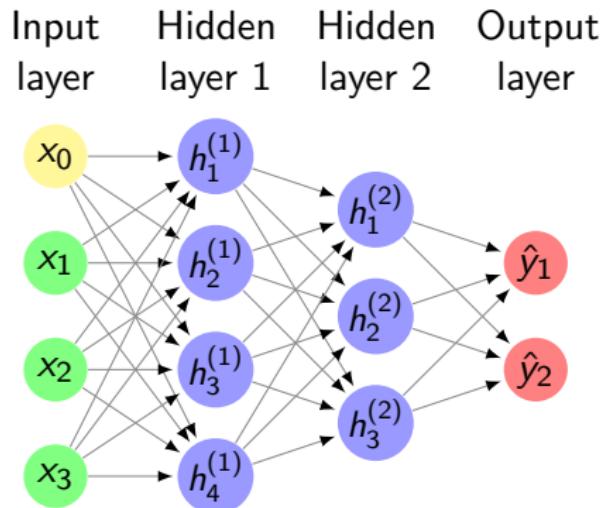


Figure: Neural Network Architecture

Key Components of Neural Networks

- **Neurons:** Fundamental processing units
 - Governed by threshold and activation function
 - **Connections:** Links between neurons
 - Transmit information between layers
 - **Weights and Biases:** Control signal strength
 - Adjusted during learning process
 - **Activation Functions:** Introduce non-linearity
 - Enable modeling of complex patterns

Single-Layer Perceptron (SLP)

- Simplest neural network (Rosenblatt, 1958)
- Solves simple problems (AND, OR, NOR gates)
- **Operation:**
 - Takes inputs $\{1, x_1, \dots, x_n\}$
 - Applies weights and bias
 - Passes through activation function
- Mathematical representation:
$$y = f(\sum_{i=1}^n w_i x_i + b_0 * 1)$$

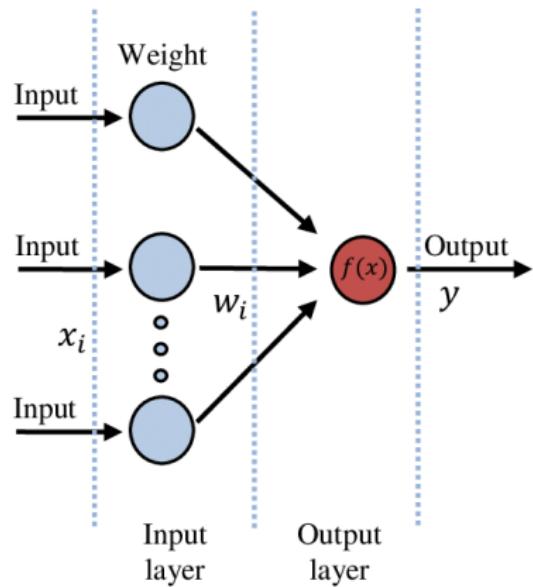


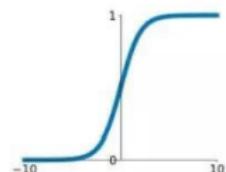
Figure: Single-layer perceptron.

Activation Functions

- **ReLU (Rectified Linear Unit):**
 - $f(x) = \max(0, x)$
 - Mitigates vanishing gradient problem
 - Most popular in modern networks
- **Tanh (Hyperbolic Tangent):**
 - $f(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$
 - Values between -1 and 1
 - Common in hidden layers
- **Sigmoid:**
 - $f(x) = \frac{1}{1+e^{-x}}$
 - Used in binary classification

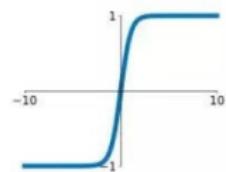
Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



tanh

$$\tanh(x)$$



ReLU

$$\max(0, x)$$

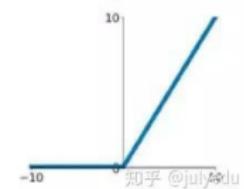


Figure: Some common activation functions.

Multi-Layer Perceptrons (MLP)

- Fully connected dense layers
- Transforms input data across dimensions
- **Components:**
 - Input layer
 - One or more hidden layers
 - Output layer
- Models complex input-output relationships
- Mathematically:
$$h_j^{(l)} = f_l(\sum_{i=1}^m w_{i,l}x_i + b_{j,l} * 1)$$

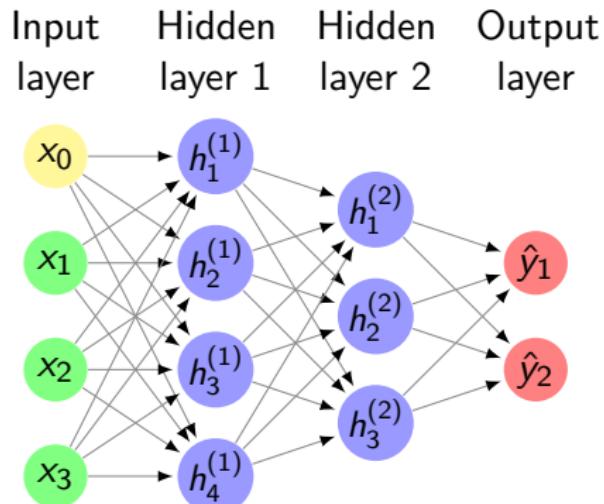


Figure: Multy-layer perceptron

Mathematical Representation of MLP

- Output of j -th neuron in l -th layer:

$$h_j^{(l)} = f_l \left(\sum_{i=1}^m w_{i,l} x_i + b_{j,l} * 1 \right) \quad (2)$$

where:

- $h_j^{(l)}$ is neuron output
 - $w_{i,l}$ are connection weights
 - $b_{j,l}$ is bias term
 - f_l is activation function
- Complete network representation:

$$y = NN(x; \theta) = W^{(l)} f^{(l)} \odot (W^{(l-1)} \odot f^{(l-1)}(\dots) + b_{l-1}) + b_l \quad (3)$$

where θ represents all parameters

Table of Contents

Theorem (Universal Approximation Theorem)

A sufficiently large neural network with at least one hidden layer can approximate any continuous function to an arbitrary degree of accuracy

1 What is a neural network?

- History
- Basic Structure & Components

2 Why at least one hidden layer? (aka Deep Learning)

3 How do we make it approximate any continuous function?

- Learning tasks & Examples
- Types of Neural Networks

What is Deep Learning?

- Subset of Machine Learning
- **Utilizes multiple processing layers**
- Enables computational models to automatically learn features and patterns
- No manual feature engineering required
- Learns hierarchical data representations

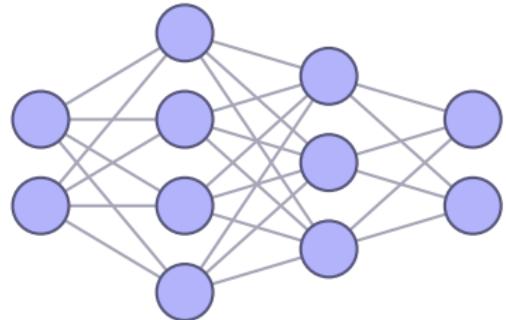
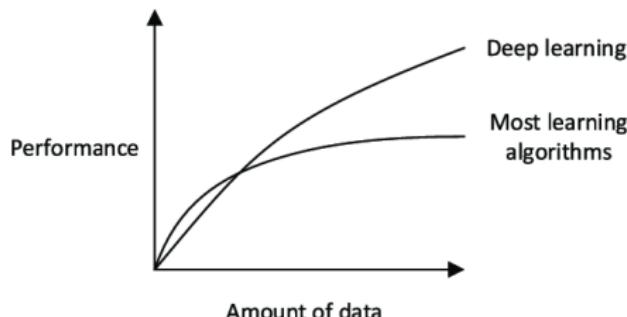


Figure: Deep Neural Network Architecture

Deep Learning vs. Machine Learning

Feature	Machine Learning	Deep Learning
Feature Extraction	Manual feature engineering	Automatic feature learning
Scalability	Struggles with large datasets	Excels with large datasets
Computation	Less computing power	Requires GPUs/high-performance
Performance (Complex Tasks)	Limited	Superior for image/speech/text



Advantages of Deep Learning

- **Handles large-scale data efficiently**
 - Performance improves with more data
- **Learns directly from raw data**
 - Eliminates manual feature engineering
- **Excels in complex tasks**
 - Image recognition
 - Speech processing
 - Autonomous systems
- **Better generalization**
- **Higher accuracy** with sufficient data

Table of Contents

Theorem (Universal Approximation Theorem)

A sufficiently large neural network with at least one hidden layer can approximate any continuous function to an arbitrary degree of accuracy

1 What is a neural network?

- History
- Basic Structure & Components

2 Why at least one hidden layer? (aka Deep Learning)

3 How do we make it approximate any continuous function?

- Learning tasks & Examples
- Types of Neural Networks

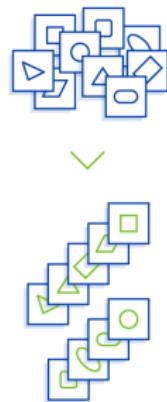
Overview of Learning Paradigms

TYPES OF MACHINE LEARNING

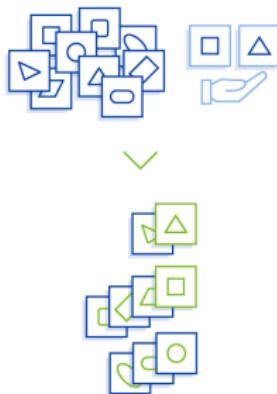
SUPERVISED
LEARNING



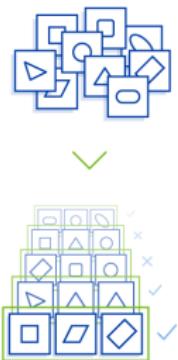
UNSUPERVISED
LEARNING



SEMI-SUPERVISED
LEARNING



REINFORCEMENT
LEARNING



Supervised Learning

- Model learns from **labeled input-output pairs**
- Predicts outputs based on inputs
- Compares predictions with known outputs
- Adjusts parameters to **minimize errors**
- **Applications:**
 - Image classification
 - Speech-to-text conversion
 - Medical diagnosis

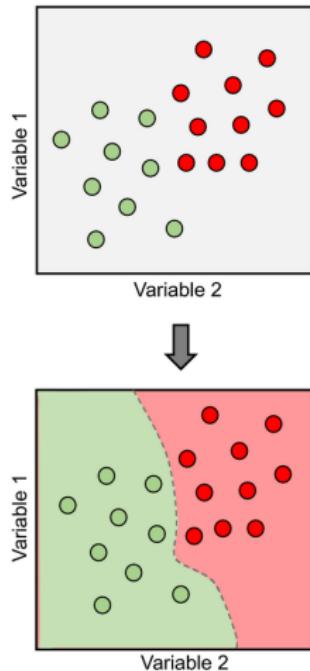


Figure: Supervised Learning

Example Supervised learning

- Plasma fusion reactor JET in Culham (UK)
- 30-second experiment generates 50 GB of data
- reconstruct the 2D plasma profile takes 1 hour per time step $\Delta t = 0.005s$
- total simulation time would be 250 days
- instead: use Convolutional Neural Network (CNN) to learn from reconstructions that have been computed at JET for all the experimental campaigns between 2011 and 2016
- multiple orders of magnitude speed up (250 days to less than 1 sec)

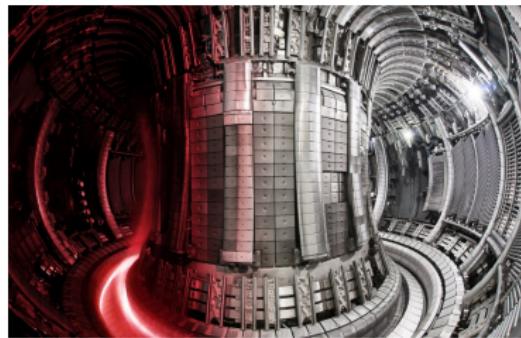


Figure: Nuclear fusion reactor JET.

Unsupervised Learning

- Learns from **unlabeled data**
- No explicit output targets
- Identifies **structures, clusters, relationships**
- **Applications:**
 - Customer segmentation
 - Anomaly detection
 - Dimensionality reduction

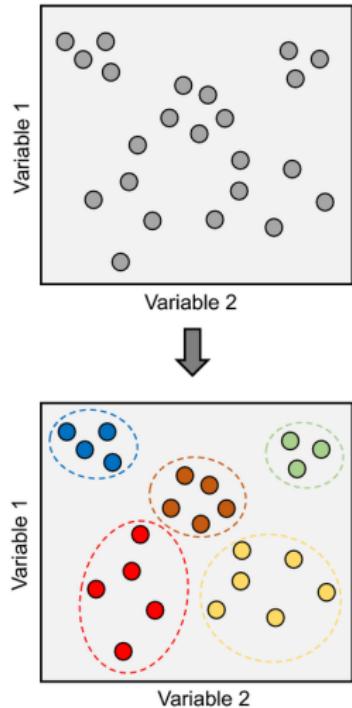


Figure: Unsupervised Learning

Example: Physics-Informed Neural Network

- No data available; only the input
- learns by a loss function that integrates the problem PDE

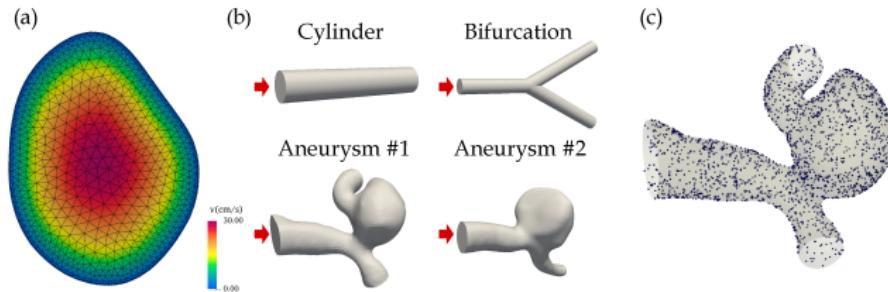
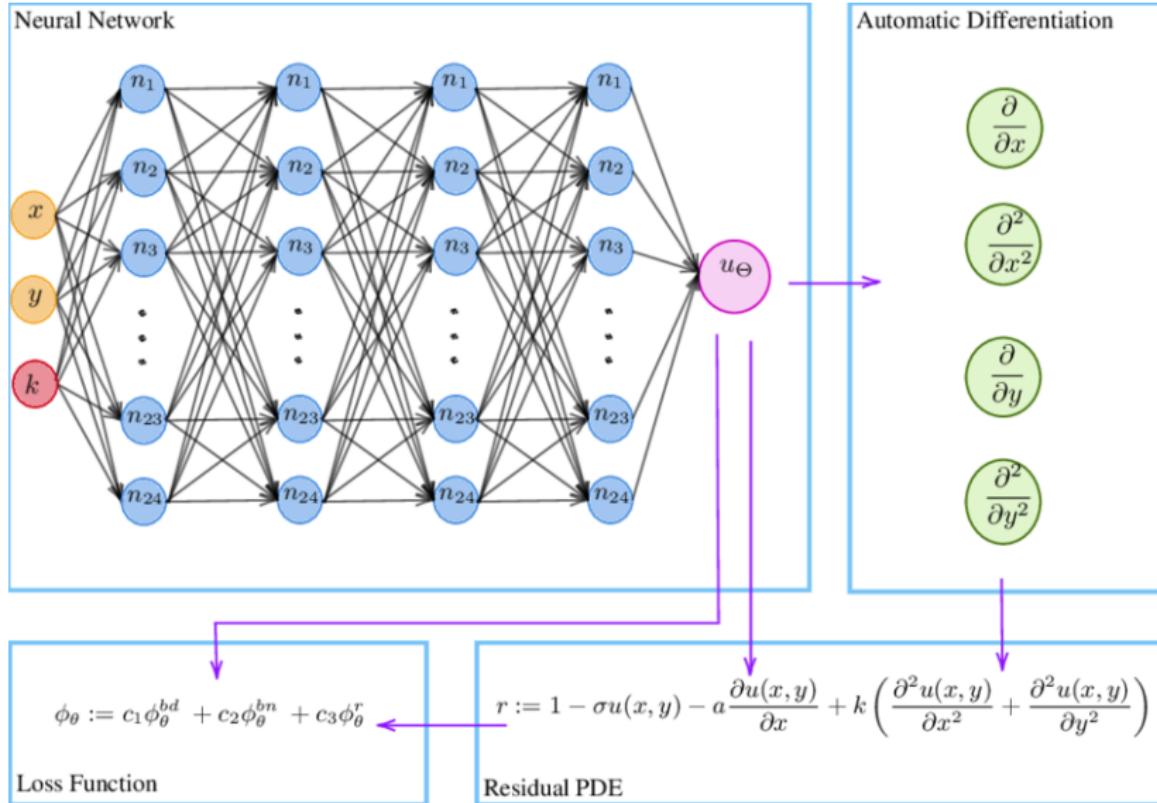


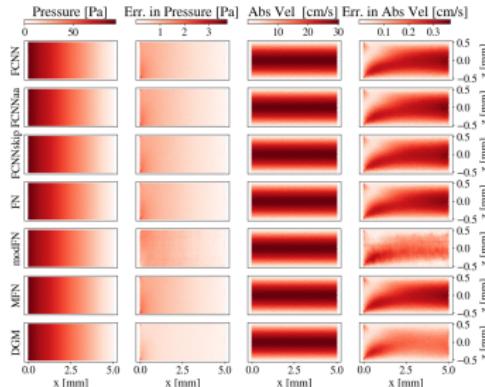
Figure: Semi-Supervised Learning

Example: Physics-Informed Neural Network

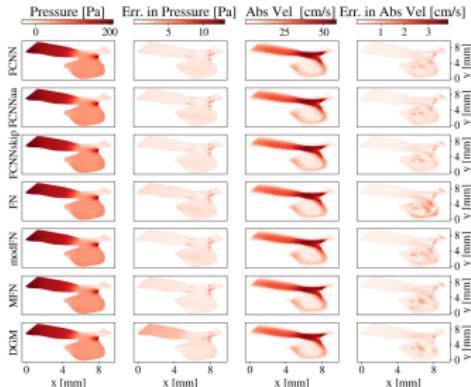


Example: Physics-Informed Neural Network

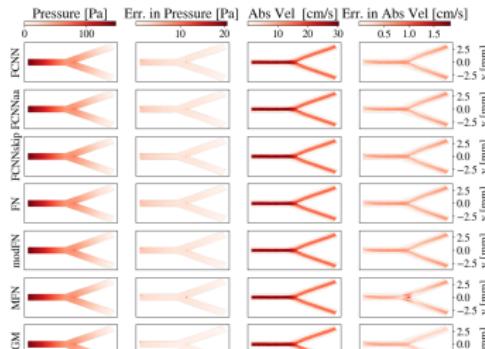
Cylinder



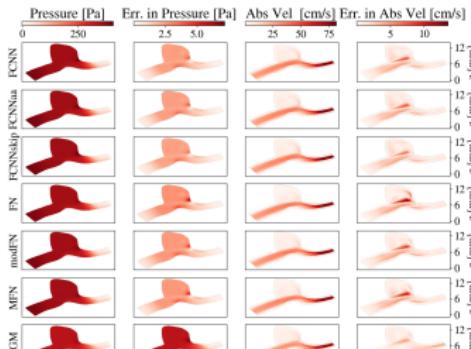
Aneurysm #1



Bifurcation



Aneurysm #2



Semi-Supervised Learning

- **Hybrid approach**
- Combines labeled and unlabeled data
- Useful when labeling is expensive
- Learns efficiently with **small amount of labeled data**
- **Applications:**
 - Speech recognition
 - Medical image analysis
 - Web page classification

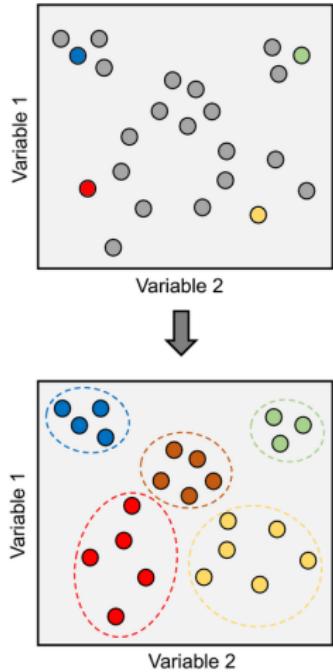


Figure: Semi-Supervised Learning

Example: Wind Farm Optimization

- Wind farm optimization takes a lot of simulations
- solution is to use a few high-resolution (HR, expensive) and many low-resolution solutions (LR, cheap)

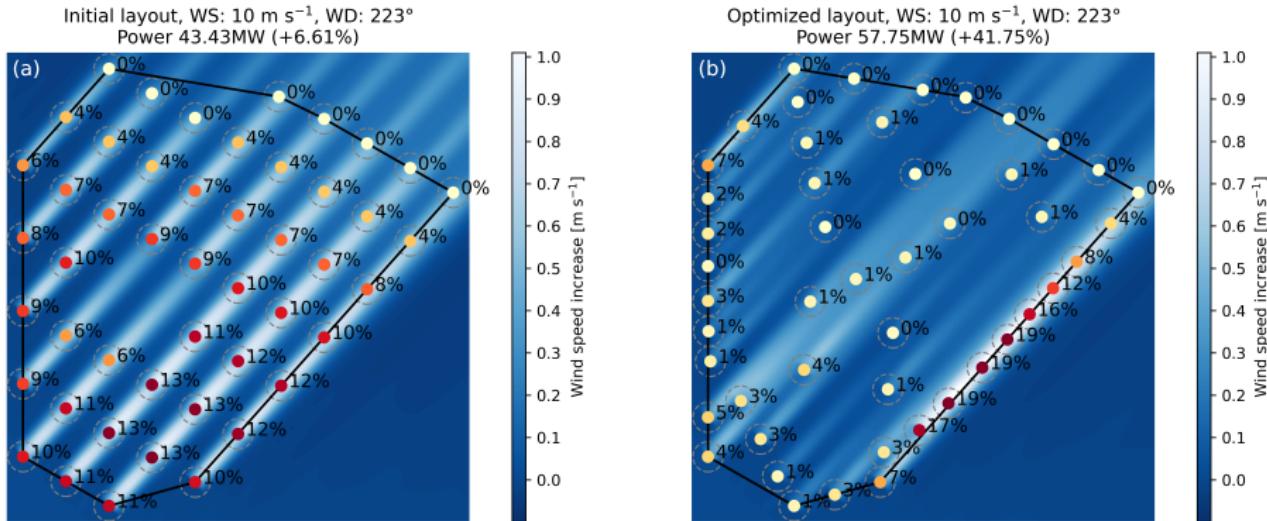


Figure: Example wind farm optimization

Example: Wind Farm Optimization

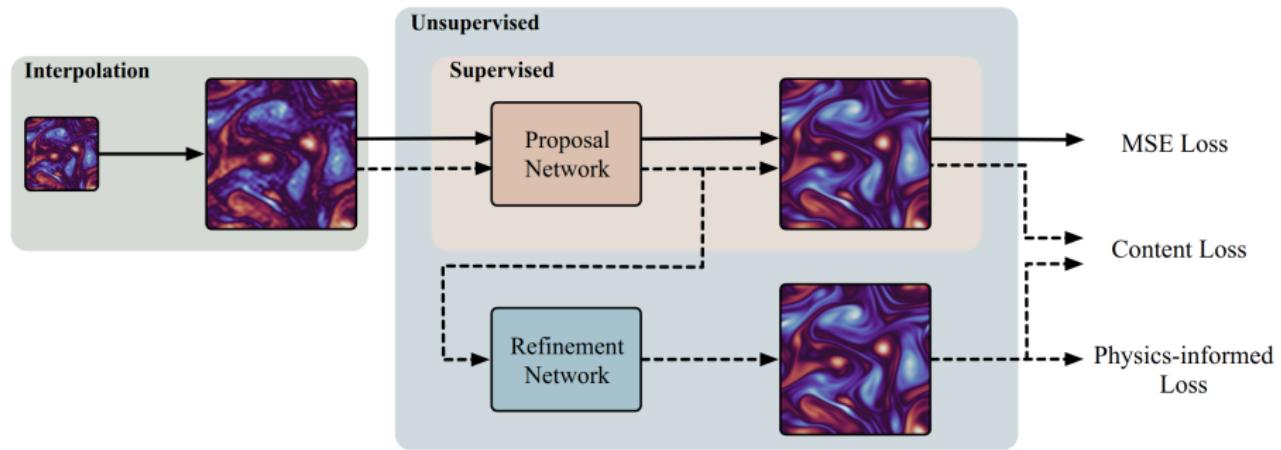


Figure: Example wind farm optimization

Reinforcement Learning

- Learning through **trial and error**
- Agent interacts with environment
- Takes actions, receives feedback
- Learns to **maximize long-term rewards**
- **Applications:**
 - Autonomous robots
 - Game AI (AlphaGo, chess)
 - Trading strategies

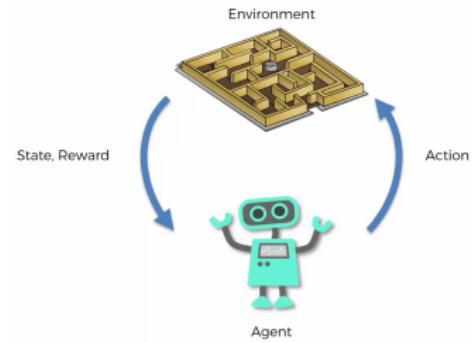
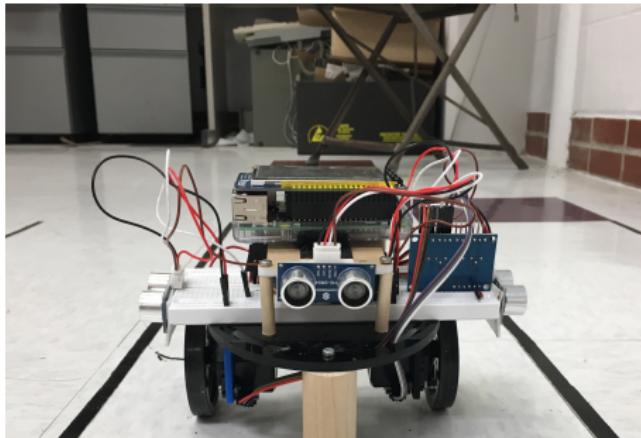


Figure: Reinforcement Learning

- autonomous robots (e.g., in a factory to drive parts)
- task: holding the lane from optical input by changing wheel movement
- full simulation would include solving the balance of linear momentum
- get a reward for holding the lane, get a penalty for leaving it
- training via interaction with the environment
- start training in the virtual environment
- move to real-life experiments



Types of Neural Networks

- **Feedforward Neural Networks (FNNs)**
 - Simplest type, one-directional information flow
 - Applications: Basic classification, simple regression
- **Multilayer Perceptron (MLP)**
 - At least one hidden layer, nonlinear activation
 - Applications: Digit recognition, stock prediction
- **Convolutional Neural Networks (CNNs)**
 - Image processing, spatial feature extraction
 - Applications: Image classification, medical imaging
- **Recurrent Neural Networks (RNNs)**
 - Sequential data, feedback loops
 - Applications: NLP, time series forecasting

Convolutional Neural Networks

- Specialized for **image processing**
- **Key components:**
 - Convolutional layers
 - Pooling layers
 - Fully connected layers
- Automatic feature extraction
- Spatial hierarchy learning

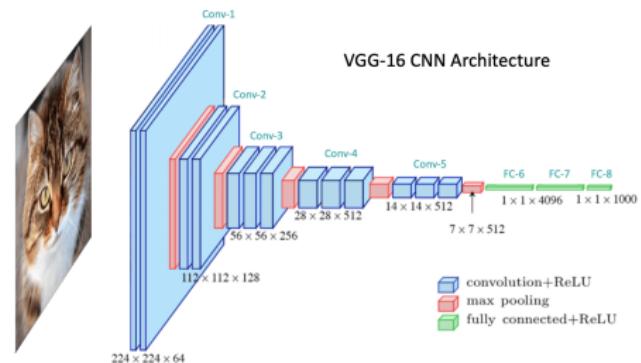


Figure: CNN Architecture

Recurrent Neural Networks

- Designed for **sequential data**
- Contains **feedback loops**
- Information persists over time
- Memory of past inputs
- Variants:
 - LSTM (Long Short-Term Memory)
 - GRU (Gated Recurrent Unit)

Recurrent Neural Network

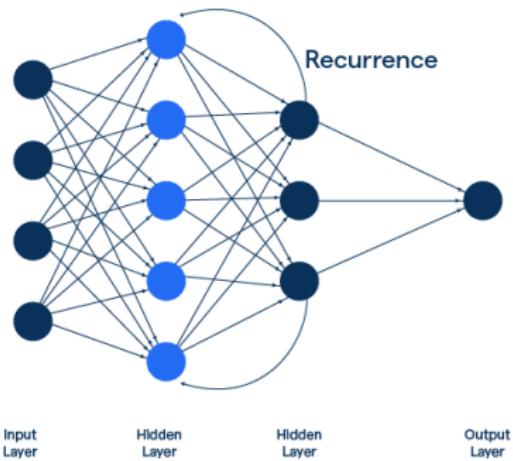


Figure: RNN Architecture

Discussion & Future Directions

- **Current challenges:**

- Interpretability and explainability
- Data efficiency and transfer learning
- Energy consumption and environmental impact

- **Emerging trends:**

- Foundation models (large pre-trained models)
- Self-supervised learning
- Multimodal learning
- Neuromorphic computing

- **Research opportunities:**

- Ethical AI and fairness
- Robust and secure deep learning
- Domain adaptation techniques

Summary

- **Deep learning** is a subset of machine learning enabling automatic feature learning
- **Four learning paradigms:** supervised, unsupervised, semi-supervised, reinforcement
- **Neural networks** are the foundation of deep learning
- **Key network types:** FNN, MLP, CNN, RNN
- **Activation functions** introduce non-linearity
- **Universal Approximation Theorem** explains theoretical power
- Deep learning continues to transform numerous industries

References I