

Introduction to Deep Learning: Part II

How to Train a Deep Neural Network

Dr. Yaohua Zang & Vincent Scholz

April 30th, 2025

Recap (I)





Figure: Neural Network Architecture


Recap (II)

TYPES OF MACHINE LEARNING

SUPERVISED LEARNING




UNSUPERVISED LEARNING




SEMI-SUPERVISED LEARNING



REINFORCEMENT LEARNING



Overview training process



Overview

1 Preparation

- Selecting Network Architecture
- Initialization of Network Parameters

2 Forward Propagation

3 Training

- Loss Function
- Backward Propagation (Optimization)

4 Summary

Table of Contents

1 Preparation

- Selecting Network Architecture
- Initialization of Network Parameters

2 Forward Propagation

3 Training

- Loss Function
- Backward Propagation (Optimization)

4 Summary

Selecting Network Architecture

The most critical aspect of deep learning that primarily influences training outcomes:

Key considerations:

- Data type and quantity available
- Learning task objective (supervised, unsupervised, reinforcement)

Architecture selection process:

- ① Look for pre-trained networks that (only) require fine-tuning
- ② Research literature for similar tasks or applications
- ③ Research wider literature (e.g., from informatics)
- ④ Only then consider making a NN yourself

Guess the number of neurons: GPT-4

Possible range: 10^0 to 10^{12}


Guess the number of neurons: GPT-4

Possible range: 10^0 to 10^{12}

Answer: $10^{11} = 100$ Billion


Guess the number of neurons: Learning a material law

Possible range: 10^0 to 10^{12}



Guess the number of neurons: Learning a material law

Possible range: 10^0 to 10^{12}



Guess the number of neurons: Learning a material law

Possible range: 10^0 to 10^{12}

Answer: 4

Architecture Considerations

When analyzing your data:

- Determine input/output dimensions
- Consider methods to integrate physical information
- Start with generous number of layers and neurons
 - Include at least one hidden layer
 - When number of parameters equals number of data pairs, exact learning becomes possible
- Select appropriate activation function (ReLU often suitable)
- Fine-tune hyperparameters iteratively

Parameter Initialization

Proper initialization ensures:

- Efficient learning
- Prevention of vanishing/exploding gradients
- Improved convergence speed

Common initialization methods:

- Zero Initialization (Not recommended)
- Constant Initialization (Rarely used)
- Uniform Initialization (Less common)
- Normal Initialization (Most common)

Advanced Initialization Techniques

- **Xavier (Glorot) Initialization**

- Weights from normal distribution: mean 0, variance $\frac{1}{n}$
- Best for sigmoid and tanh activation functions
- Maintains variance of activations across layers

- **He Initialization**

- Weights from normal distribution: mean 0, variance $\frac{2}{n}$
- Recommended for ReLU and Leaky ReLU
- Prevents vanishing gradient problem

- **Orthogonal Initialization**

- Uses orthogonal matrix
- Preserves gradient norms during backpropagation
- Useful in deep architectures for numerical stability

Table of Contents

1 Preparation

- Selecting Network Architecture
- Initialization of Network Parameters

2 Forward Propagation

3 Training

- Loss Function
- Backward Propagation (Optimization)

4 Summary

Forward Propagation

The process of passing input data through network layers to compute output:

$$\hat{y} = NN(x; \theta) = W^{(I)} f^{(I)} \odot (W^{(I-1)} \odot f^{(I-1)}(\dots) + b_{I-1}) + b_I$$

Where:

- \hat{y} - predicted output
- θ - all trainable parameters: $\{W^{(I)}, b^{(I)}, \dots, W^{(1)}, b^{(1)}\}$
- $W^{(I)}, b^{(I)}$ - weights and biases of layer I
- $f^{(I)}$ - activation function at layer I
- \odot - element-wise multiplication




Figure: Neural Network Architecture

Forward Propagation Steps

- ① **Input Layer:** Feed input data x into network
- ② **Hidden Layers:** Each hidden layer applies:

$$z^{(l)} = W^{(l)} h^{(l-1)} + b^{(l)}$$

$$h^{(l)} = f^{(l)}(z^{(l)})$$

where $h^{(l)}$ is the activation at layer l

- ③ **Output Layer:** Final layer produces predicted output \hat{y}
- ④ **Loss Calculation:** Compare prediction to actual target using loss function

Table of Contents

1 Preparation

- Selecting Network Architecture
- Initialization of Network Parameters

2 Forward Propagation

3 Training

- Loss Function
- Backward Propagation (Optimization)

4 Summary

Loss Function (Cost Function)

Quantifies how well the neural network performs:

- ① Pass each training example (x_i, y_i) through network to get prediction
 $\hat{y}_i = NN(x_i; \theta)$
- ② Compare \hat{y}_i with actual label y_i using loss function
- ③ Compute average loss over all training examples, e.g.:


$$\mathcal{L}(\theta) = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

Regression vs Classification

The choice of loss function depends on the problem type:


- **Regression:** Predict continuous values
- **Classification:** Predict discrete labels

Classification Groups observations into "classes"



Here, the line classifies the observations into X's and O's

Regression predicts a numeric value



Here, the fitted line provides a predicted output, if we give it an input

Regression Loss Functions

① Mean Squared Error (MSE or L2 Loss):

$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

- Amplifies larger deviations (sensitive to outliers)
- Differentiable for gradient-based optimization
- Ideal for normal (Gaussian) distribution targets

② Mean Absolute Error (MAE or L1 Loss):

$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}_i|$$


- Less sensitive to outliers than MSE
- Not differentiable at $y_i = \hat{y}_i$

More Regression Loss Functions

③ Huber Loss (Smooth L1 Loss):

$$\mathcal{L}_\delta(y_i, \hat{y}_i) = \begin{cases} \frac{1}{2}(y_i - \hat{y}_i)^2, & |y_i - \hat{y}_i| \leq \delta \\ \delta \cdot (|y_i - \hat{y}_i| - \frac{1}{2}\delta), & \text{otherwise} \end{cases}$$

- Combines advantages of MSE and MAE
- Quadratic for small errors, linear for large errors
- Robust to outliers while maintaining differentiability
- Hyperparameter δ determines transition point



Classification Loss Functions

① Binary Cross-Entropy (Log Loss):

$$\mathcal{L} = -\frac{1}{N} \sum_{i=1}^N (y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i))$$

- For binary classification tasks
- Penalizes incorrect predictions, encourages high confidence
- Works well for binary labels (e.g., spam vs. not spam)
- Logarithmic penalty heavily punishes confident incorrect predictions

	$y = 0$	$y = 1$
$\hat{y} = 0.01$	0.01	4.6
$\hat{y} = 0.99$	4.6	0.01

Multi-Class Classification Loss

Table: One hot encoding for Multi-Class Classification

Feature name	# Feature	One-hot encoding
Cat	1	[1, 0, 0]
Bee	2	[0, 1, 0]
Dog	3	[0, 0, 1]

Multi-Class Classification Loss

Categorical Cross-Entropy (Softmax Loss):

$$\mathcal{L}(y_i, \hat{y}_i) = - \sum_j^c t_j \log\left(\frac{e^{s_p}}{\sum_j^c e^{s_j}}\right) = - \log\left(\frac{e^{s_p}}{\sum_j^c e^{s_j}}\right)$$

Where:

- $y_i = (t_1, \dots, t_c)$, $\hat{y}_i = (s_1, \dots, s_c)$
- C = number of classes
- s_j = predicted score for class j
- s_p = score of correct class
- One-hot encoded targets: $t_i = 0, \forall i \neq p$ and $t_p = 1$

Summary of Loss Functions

Task Type	Loss Functions	Characteristics
Regression	MSE (L2)	<ul style="list-style-type: none">- Smooth,- differentiable,- sensitive to outliers
	MAE (L1)	<ul style="list-style-type: none">- Robust to outliers,- non-differentiable at zero
	Huber	<ul style="list-style-type: none">- Balanced approach between MSE and MAE
Binary Classification	Binary	<ul style="list-style-type: none">- Encourages confident predictions
Multi-Class Classification	Categorical	<ul style="list-style-type: none">- softmax for probability distribution
	Cross-Entropy	<ul style="list-style-type: none">- Penalizes incorrect predictions
	Cross-Entropy	<ul style="list-style-type: none">- Not suitable for multi-label

Backward Propagation

Backward Propagation of Errors:

- Fundamental algorithm for training neural networks
- Minimizes discrepancy between predicted and actual outputs
- Adjusts weights and biases using gradient of loss function

Training objective:

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta) = \frac{1}{N} \sum_{i=1}^N (y_i - NN(x_i; \theta))^2$$

Understanding Backpropagation

Key aspects:

- **Iterative process** minimizing cost function
- Updates model parameters in direction of negative gradient
- Uses **chain rule** to propagate error backward through network
- Computes partial derivatives efficiently layer by layer
- Updates parameters using gradient descent or variants

Picture Backpropagation




Figure: Backpropagation

Mathematical Formulation

For a single neuron:

$$z = w_1x_1 + w_2x_2 + \cdots + w_nx_n + b$$
$$a = f(z)$$

Using the chain rule to compute gradients:

$$\frac{\partial \mathcal{L}}{\partial w_i} = \frac{\partial \mathcal{L}}{\partial a} \cdot \frac{\partial a}{\partial z} \cdot \frac{\partial z}{\partial w_i}$$
$$\frac{\partial \mathcal{L}}{\partial b} = \frac{\partial \mathcal{L}}{\partial a} \cdot \frac{\partial a}{\partial z} \cdot \frac{\partial z}{\partial b}$$

Update weights and biases using gradient descent:

$$w_i \leftarrow w_i - \eta \frac{\partial \mathcal{L}}{\partial w_i}$$
$$b \leftarrow b - \eta \frac{\partial \mathcal{L}}{\partial b}$$

where η is the learning rate.

Gradient Descent

Gradient Descent is a fundamental optimization algorithm that:

- Minimizes cost function by iteratively adjusting parameters
- Moves in direction of negative gradient to find minimum

Parameter update rule:

$$\theta_{t+1} = \theta_t - \eta \frac{\partial \mathcal{L}(\theta_t)}{\partial \theta_t}$$

Where:

- θ_t = model parameters at iteration t
- η = learning rate (controls step size)
- $\frac{\partial \mathcal{L}(\theta_t)}{\partial \theta_t}$ = gradient of loss function

Gradient Descent Picture




Figure: Gradient Descent.

Importance of Learning Rate


Learning rate η is a crucial tuning parameter:

- Determines step size at each iteration
- Balances optimization time and accuracy
- **Small η :** Small steps, slow convergence
- **Large η :** Risk of overshooting minimum


The step size depends on:

- Gradient of cost function (steepness of slope)
- Learning rate value
- Training stops when gradient becomes zero

Importance of Learning Rate



Importance of Learning Rate



Variants of Gradient Descent (1)

1. Stochastic Gradient Descent (SGD)

- Updates using single example at a time
- Suitable for large datasets

$$\theta_{t+1} = \theta_t - \eta \frac{\partial \mathcal{L}_i(\theta_t)}{\partial \theta_t}, \quad \text{where } \mathcal{L}_i(\theta_t) = (y_i - NN(x_i; \theta))^2$$

Pros:

- Faster updates, suitable for large datasets
- Can escape local minima due to noisy updates

Cons:

- High variance in updates, may lead to instability

Variants of Gradient Descent (2)

2. Mini-Batch Gradient Descent

- Updates using a small batch of examples
- Balance between computation time and precision


$$\theta_{t+1} = \theta_t - \eta \frac{\partial \mathcal{L}_b(\theta_t)}{\partial \theta_t}, \quad \text{where } \mathcal{L}_b(\theta_t) = \frac{1}{m} \sum_{i=1}^m (y_i - NN(x_i; \theta))^2$$

where $m \ll n$ is the mini-batch size.

Pros:

- Reduces variance while maintaining computational efficiency

Comparison SGD / Minibatch



Variants of Gradient Descent (3)

3. Momentum-Based Gradient Descent

- Adds information from preceding steps to next step
- Helps penetrate flat areas and noisy gradients

$$\theta_{t+1} = \theta_t - \eta \frac{\partial \mathcal{L}_i(\theta_t)}{\partial \theta_t} + \alpha \Delta \theta_t$$

Where:

- $\Delta \theta_t$ = weight change at previous step
- $\alpha \in (0, 1)$ = momentum coefficient (typically 0.9)

Pros:

- Helps escape local minima
- Reduces oscillations and speeds up convergence

Comparison SGD / Momentum-based




Table of Contents

1 Preparation

- Selecting Network Architecture
- Initialization of Network Parameters

2 Forward Propagation

3 Training

- Loss Function
- Backward Propagation (Optimization)

4 Summary

Training Deep Neural Networks: Summary

① Network Architecture

- Key to data learning and processing
- Based on data quality/quantity and domain information

② Parameter Initialization

- Initialize with suitable techniques (Normal, Xavier, etc.)

③ Forward Propagation

- Pass input data through network layers
- Apply linear transformations and activation functions

④ Loss Function

- Measure prediction accuracy (MSE, Cross-Entropy, etc.)

⑤ Backpropagation (Optimization)

- Calculate gradients and update parameters
- Use gradient descent: $\theta_{t+1} = \theta_t - \eta \frac{\partial \mathcal{L}(\theta_t)}{\partial \theta_t}$

References & Additional Resources

- Goodfellow, I., Bengio, Y., & Courville, A. (2016). Deep learning. MIT press.
- LeCun, Y., Bengio, Y., & Hinton, G. (2015). Deep learning. *Nature*, 521(7553), 436-444.
- Rumelhart, D. E., Hinton, G. E., & Williams, R. J. (1986). Learning representations by back-propagating errors. *Nature*, 323(6088), 533-536.
- Kingma, D. P., & Ba, J. (2014). Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*.

Online Resources:

- PyTorch Documentation:
<https://pytorch.org/docs/stable/index.html>